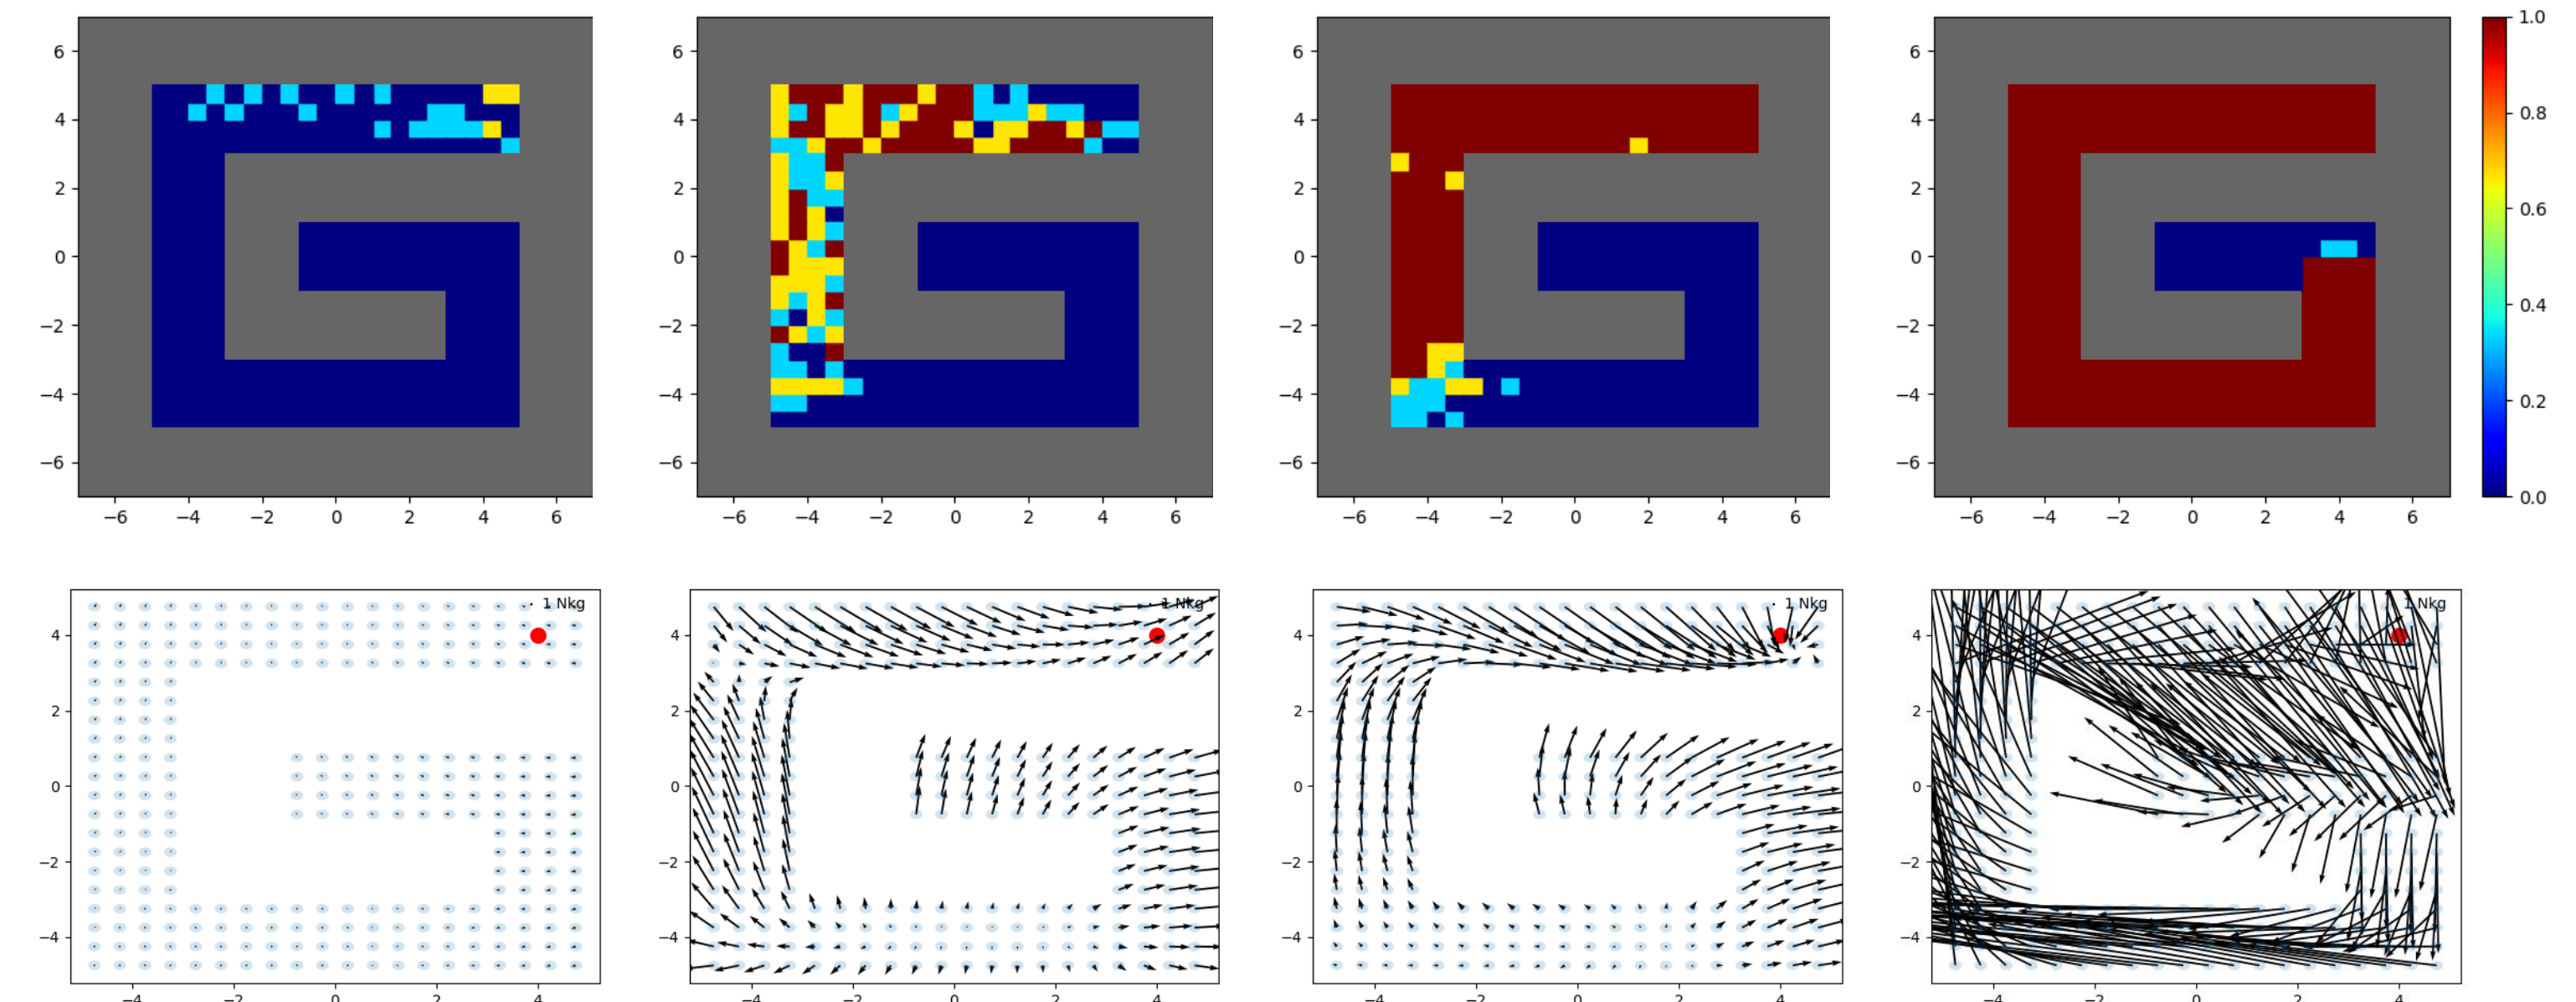


Reverse Curriculum

Amadeus Hovekamp, Megan Klaiber, Hans Nübel, Rabea Turon

Project Idea

- Reproduce the results of „Reverse Curriculum Generation for Reinforcement Learning“ [1]
 - Using Code provided by the authors of the paper
 - Using a RL environment from the paper and OpenAI baselines RL algorithm
- Reverse Curriculum idea: Start training from states close to the goal and gradually move further away



Adaptive Curriculum Generation

- Alternative to uniform sampling over all states
- Sample from Start States of current iteration
- Generate new Start States with Brownian Motion
- Select starts according to sampling method
 - From all Previous
 - Good Starts only

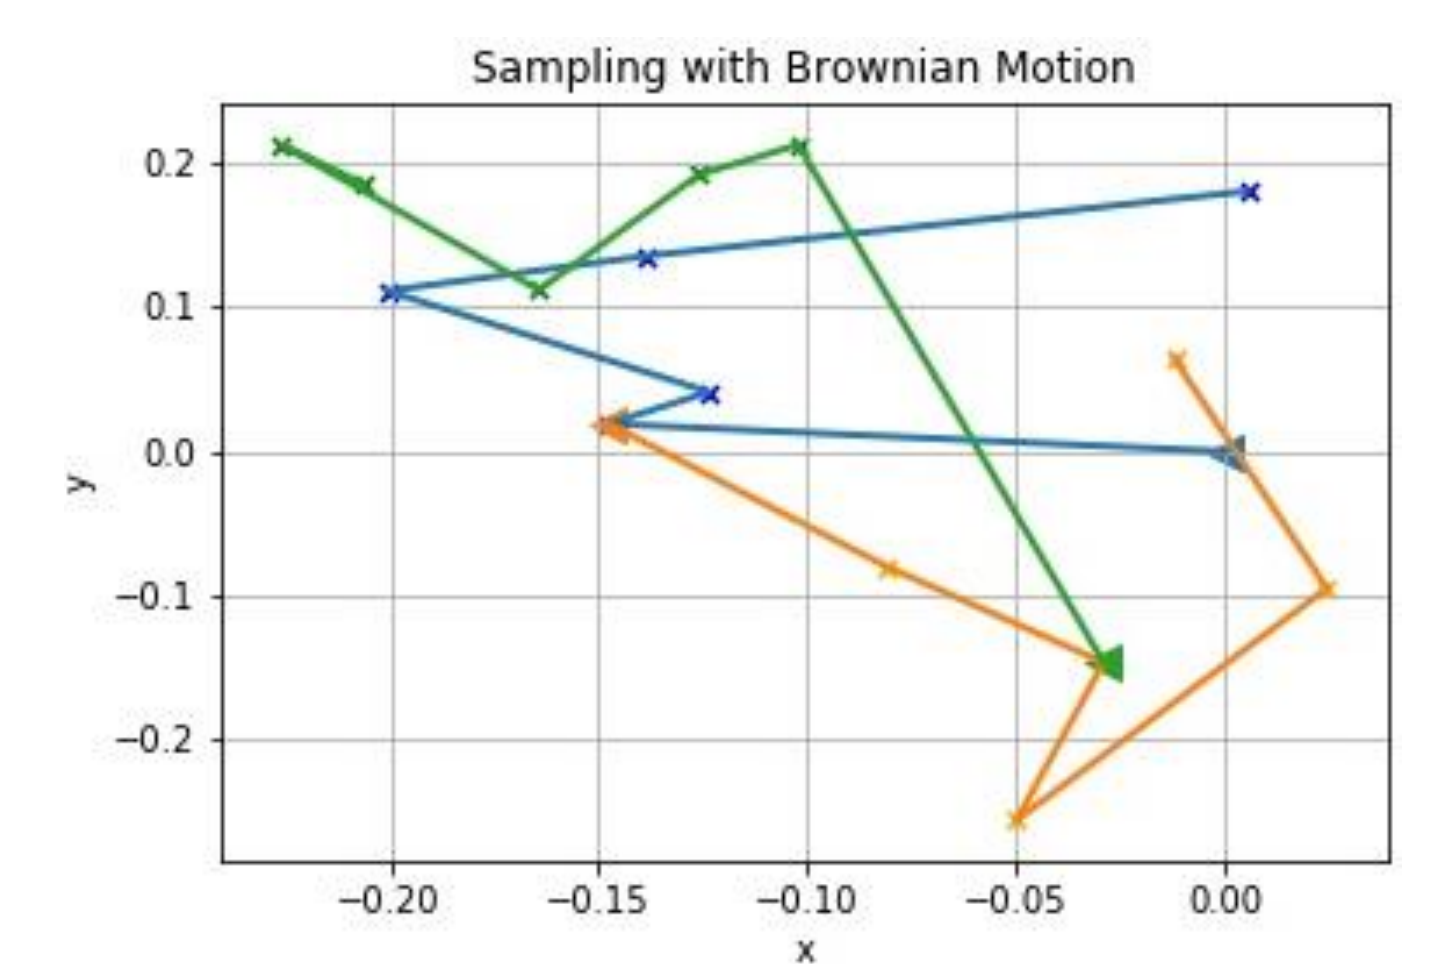
Algorithm 1: Policy Training

Input : $\pi_0, s^g, \rho_0, N_{\text{new}}, N_{\text{old}}, R_{\text{min}}, R_{\text{max}}, \text{Iter}$
Output: Policy π_N

```

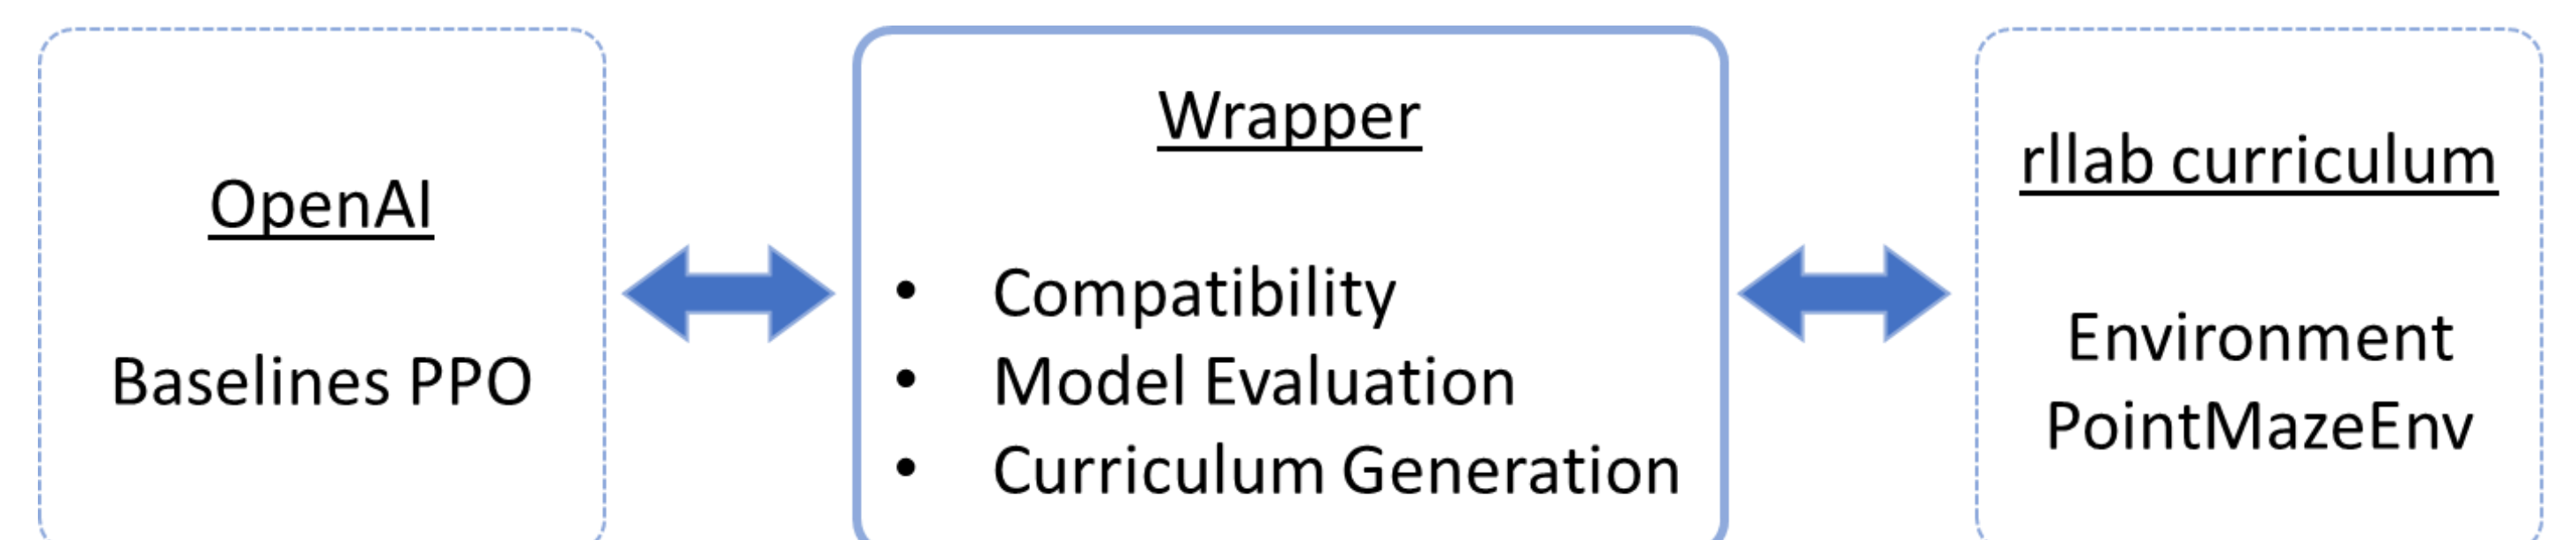
 $starts_{\text{old}} \leftarrow [s^g];$ 
 $starts, rews \leftarrow [s^g], [1];$ 
for  $i \leftarrow 1$  to  $\text{Iter}$  do
   $starts \leftarrow \text{SampleNearby}(starts, N_{\text{new}});$ 
   $starts.append(\text{sample}(starts_{\text{old}}, N_{\text{old}}));$ 
   $\rho_i \leftarrow \text{Unif}(starts);$ 
   $\pi_i, rews \leftarrow \text{train\_pol}(\rho_i, \pi_{i-1});$ 
   $starts \leftarrow \text{select}(starts, rews, R_{\text{min}}, R_{\text{max}});$ 
   $starts_{\text{old}}.append(starts);$ 
end
  
```

Algorithm 1 from [1]

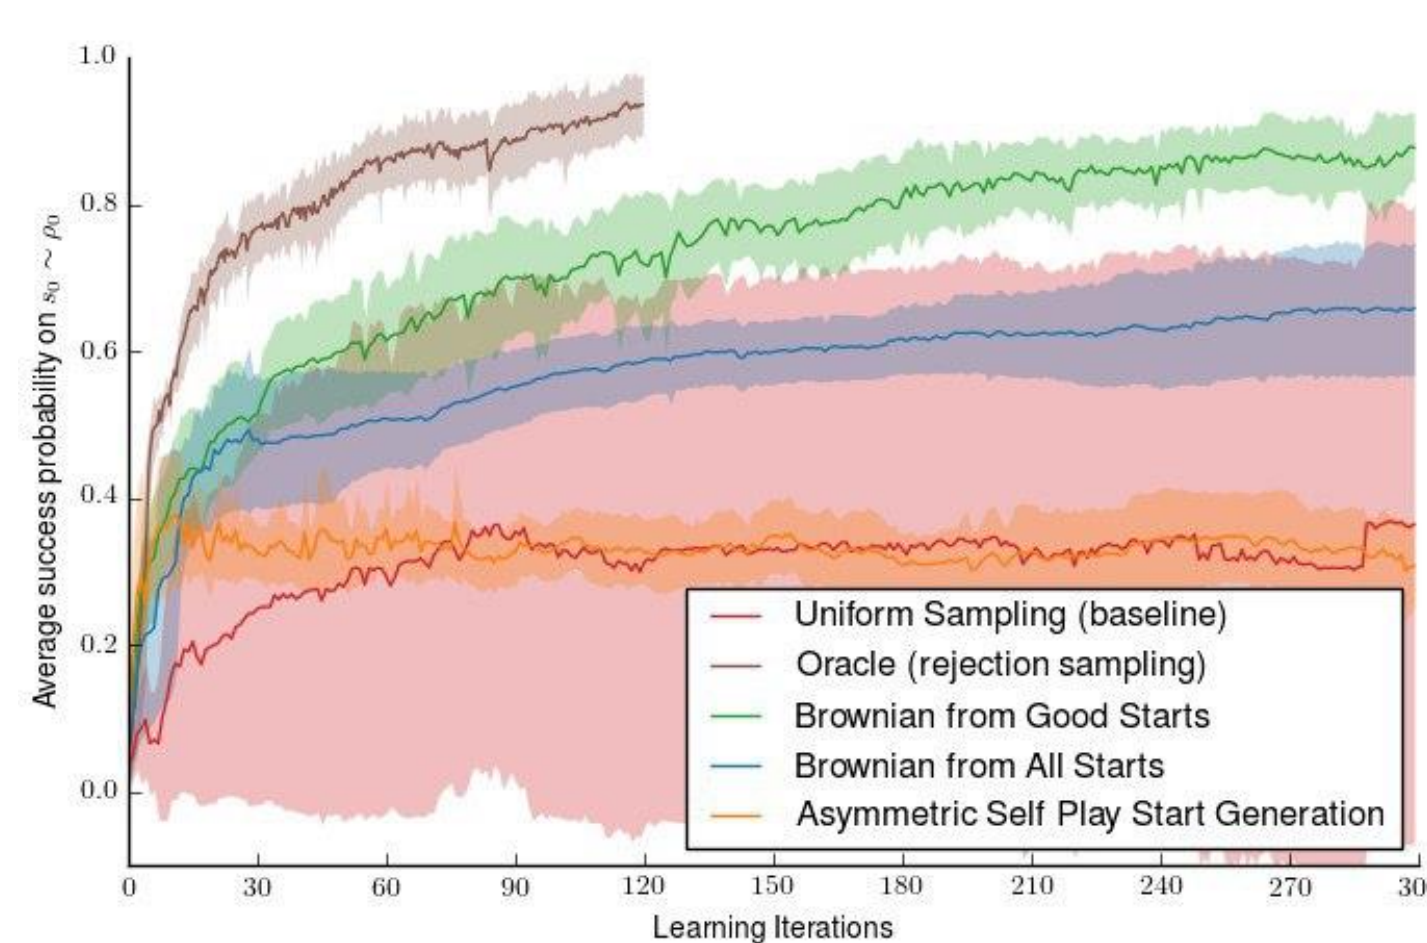


Wrapper

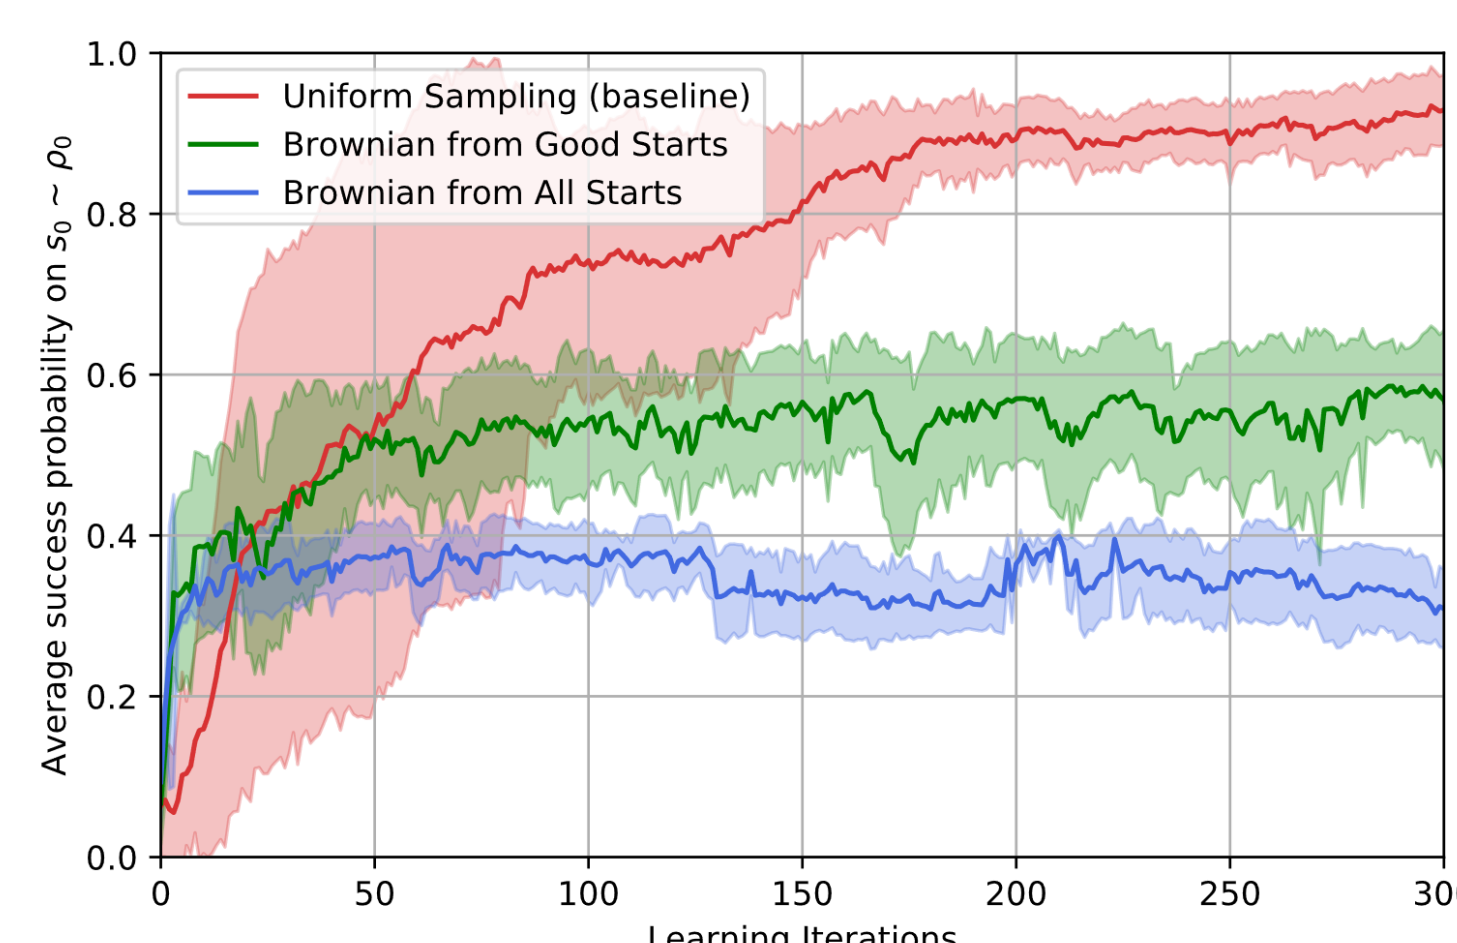
- Contains our implementation
- Connects OpenAI to rllab curriculum environment



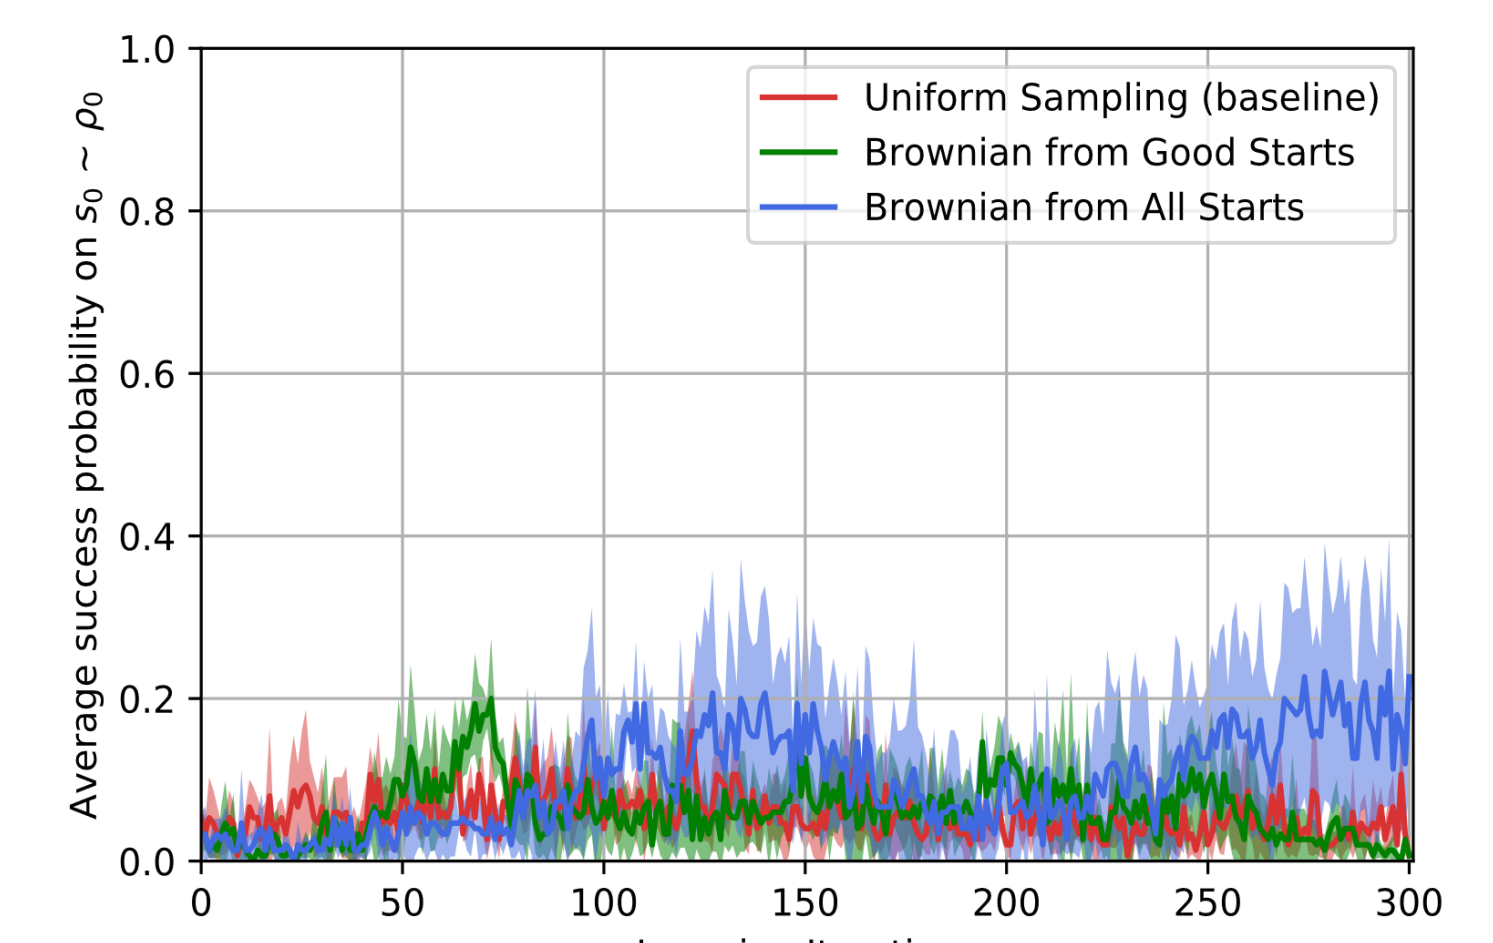
Results



(a) Original results from [1]



(b) Reproduced results with rllab curriculum code



(c) OpenAI Baselines PPO

- We were not able to reproduce the results: substantial differences between their results in (a) and the results rerunning their code (b)
- Surprisingly, uniform sampling outperformed the more sophisticated methods
- Reimplementation of Algorithm 1 and using Proximal Policy Optimization as training method shows bad results, see (c)
- Brownian Motion from both Good and All Starts performed better than Uniform Sampling in our implementation

Issues

- Problems in original code (rllab curriculum)
- Incompatibility of OpenAI algorithms and rllab curriculum environments
- Differences between code and paper

References

- [1] Florensa, C., Held, D., Wulfmeier, M., Zhang, M. & Abbeel, P.. (2017). Reverse Curriculum Generation for Reinforcement Learning. Proceedings of the 1st Annual Conference on Robot Learning, in PMLR 78:482-495